


## Optimizing Deep Learning Performance:

A Hybrid CPU-GPU Framework with Multithreading, SIMD,  
and Evaluation of Efficiency Metrics



**01** Multithreading

**02** Hardware

**03** Benchmarks

**04** CUDA tuning

**05** Outlook

Multithreading has been implemented to optimize the performance of these functions:

*tf.h:*

- matrix **\*add**(matrix \*a, matrix \*b);
- matrix **\*\*biasing**(matrix \*\*a, **int** len, matrix \*b);
- matrix **\*\*conv2d**(matrix \*a, matrix \*\*b, **int** len);
- matrix **\*flatten**(matrix \*\*a, **int** len);
- matrix **\*\*flip\_kernels**(matrix \*\*a, **int** len);
- matrix **\*\*hyperbolic\_tangent**(matrix \*\*a, **int** len);
- matrix **\*matmul**(matrix \*a, matrix \*b);
- matrix **\*\*maxpool**(matrix \*\*a, **int** len);
- matrix **\*\*relu**(matrix \*\*a, **int** len);
- matrix **\*transpose**(matrix \*a);



*mt.h:*

- **void add\_mt**(mt\_arg \*mt);
- **void biasing\_mt**(mt\_arg \*mt);
- **void conv2d\_mt**(mt\_arg \*mt);
- **void flatten\_mt**(mt\_arg \*mt);
- **void flip\_kernels\_mt**(mt\_arg \*mt);
- **void hyperbolic\_tangent\_mt**(mt\_arg \*mt);
- **void matmul\_mt**(mt\_arg \*mt);
- **void maxpool\_mt**(mt\_arg \*mt);
- **void relu\_mt**(mt\_arg \*mt);
- **void transpose\_mt**(mt\_arg \*mt);

### 1. We have implemented the following struct to be passed to the pthreads:

*mt.h:*

```
typedef struct mt_arg {  
    long idx;  
  
    matrix **a_ptr;  
  
    matrix *a;  
  
    matrix **b_ptr;  
  
    matrix *b;  
  
    int len;  
  
    matrix **c_ptr;  
  
    matrix *c;  
  
    int m;  
  
    void (*start_routine)(struct mt_arg *mt);  
} mt_arg;
```

### 2. We have developed a proof of concept:

- Each function independently created and joined its own pthreads
- The implementation was functional
- However, performance was **suboptimal**

### 3. We have implemented a thread pool:

- Thanks to the suggestion of Philipp Holzinger :)
- **Threads are now:**
  - Created once at the start using the function:
    - `void create_mt(long threads);`
  - Joined once at the end using the function:
    - `void join_mt();`
- **Implementation:**
  - Utilizes a GAsyncQueue for task management
  - The threads execute the following function:
    - `static void *start_mt(void *arg);`

### – Thread behavior:

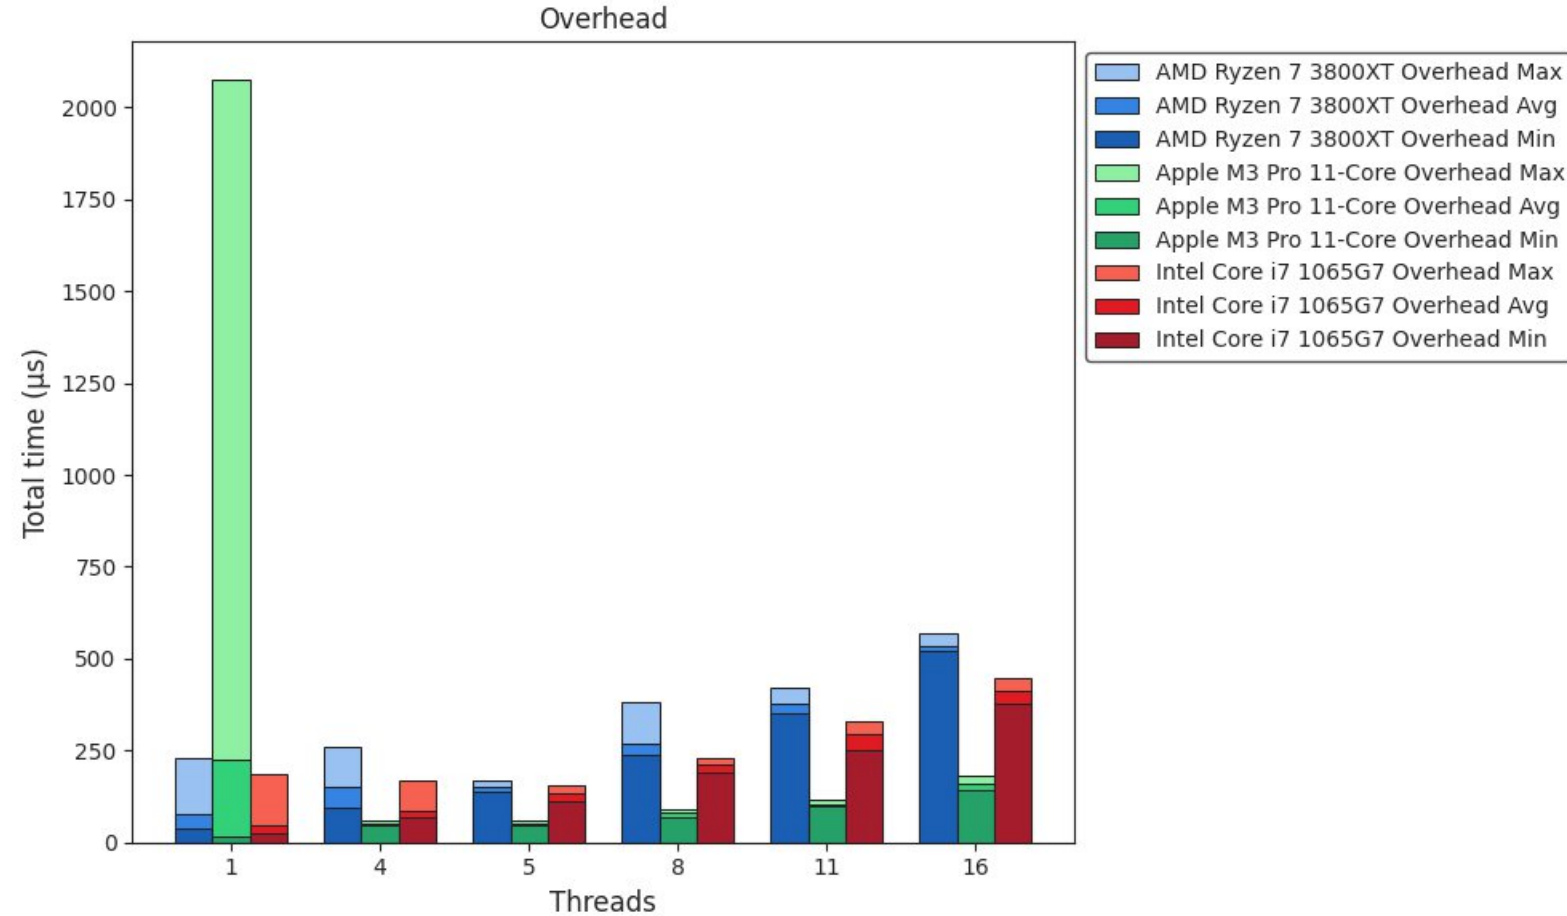
- The threads wait for elements to be added to the queue
- Upon receiving an element, they execute the provided function with its associated parameters
- This process repeats indefinitely
- Elements can be added to the queue with:
  - `void push_mt(mt_arg *mt);`
- `void join_mt();` adds a termination signal to the queue by pushing the function:
  - `static void stop_mt(mt_arg *mt);`
- This ensures that all threads exit cleanly and can be joined properly

- **A race condition was identified and debugging it required significant effort:**
  - To address this, a synchronization barrier was implemented using the function:
    - `void wait_mt();`
  - This implementation required an understanding of `pthread_cond_t` and `pthread_mutex_t`
- **Testing improvements:**
  - The existing dataset (MNIST) was determined to be too small for thorough testing
  - A larger test set was created by applying a scale factor of 20, resulting in image dimensions of  $(30 \times 20)^2$
- **Concurrency research and evaluation:**
  - Various concurrency solutions were analyzed for suitability
  - **Fork:**
    - Duplicates the calling process
    - Offers low flexibility and incurs high overhead
  - **Pthreads:**
    - The current solution being used
    - Provides full control over concurrency
  - **OpenMP:**
    - A compiler-based concurrency implementation
    - Offers lower control
    - We plan to explore its implementation in the upcoming weeks

CPU	Release dates	TDP (W)	Number of (performance) cores	Number of threads
AMD Ryzen 7 3800XT	7. Juli 2020	105	8	16
Apple M3 Pro 11-Core	30. Oktober 2023	27	5	11
Intel Core i7 1065G7	1. Juni 2019	15	4	8

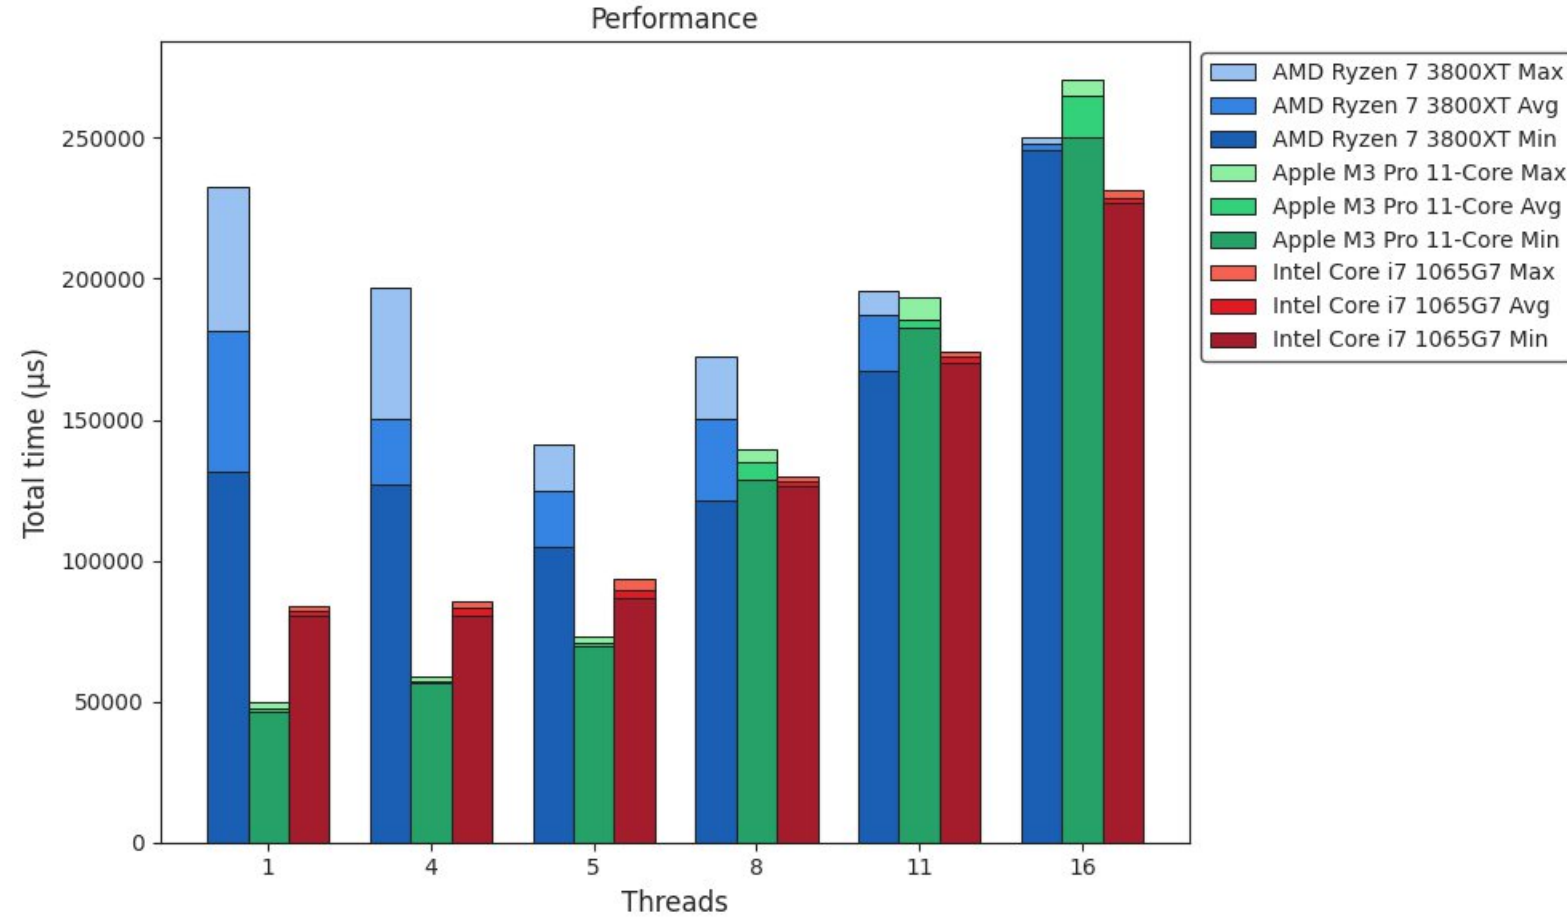
- 
- **Batch size:** 1
  - **Epochs:** 128
  - **Overhead:**
    - In microseconds ( $\mu$ s)
    - Averaged over 10 runs
    - **Sum of:**
      - `void create_mt(long threads);`
      - `void join_mt();`
  - **Performance:**
    - In microseconds ( $\mu$ s)
    - Averaged over 10 runs
    - Total time
  - **XL:**
    - In microseconds ( $\mu$ s)
    - Averaged over 10 runs
    - Image dimensions of  $(30 \times 20)^2$

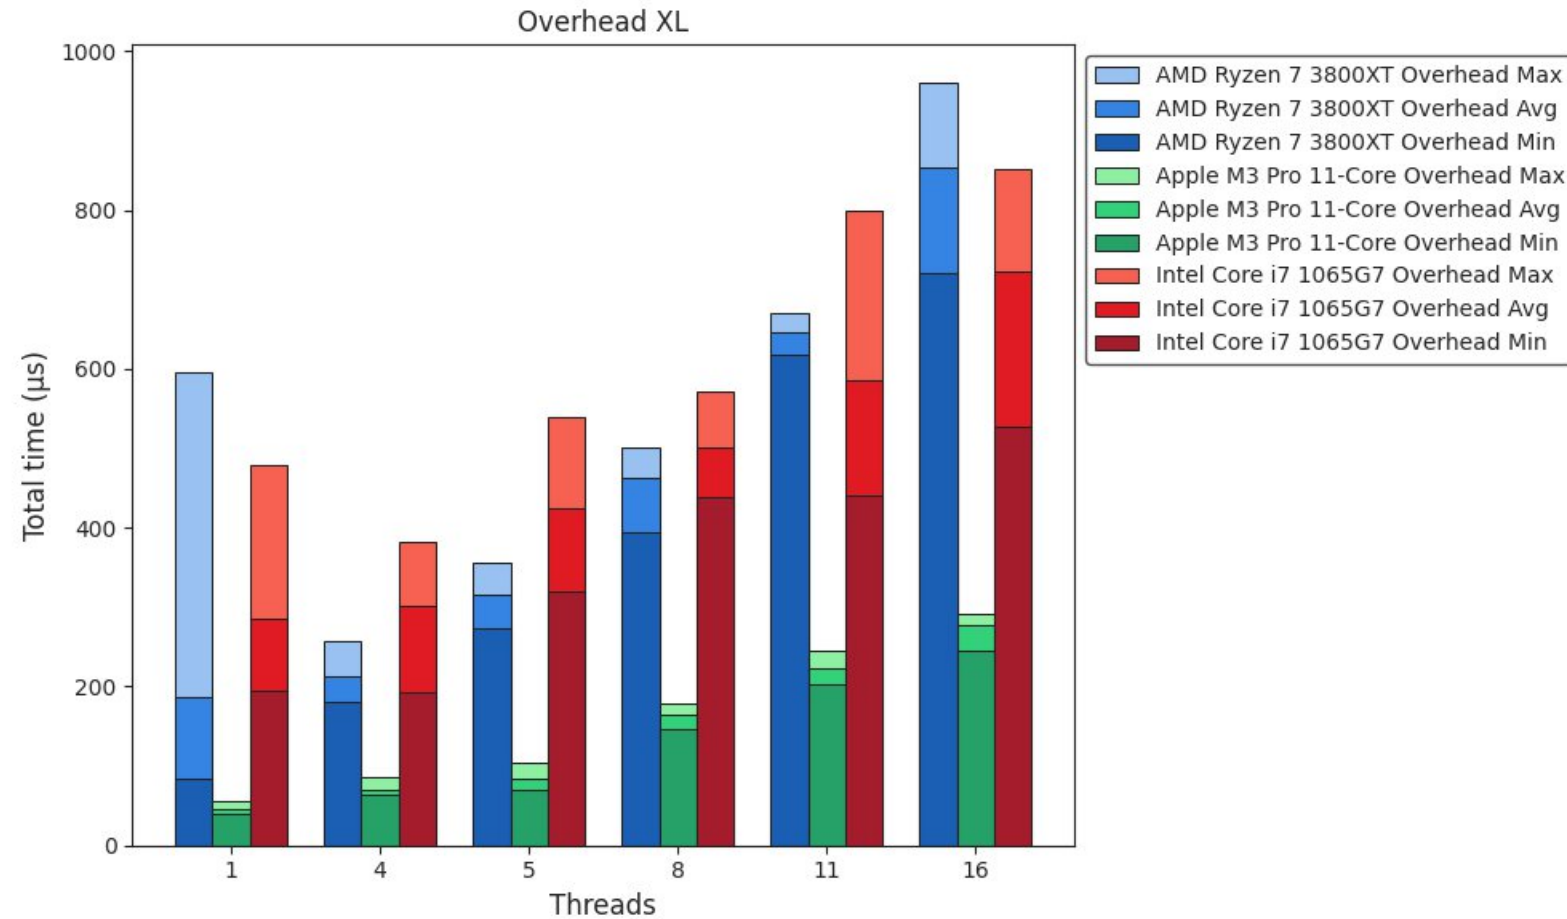




# Benchmark

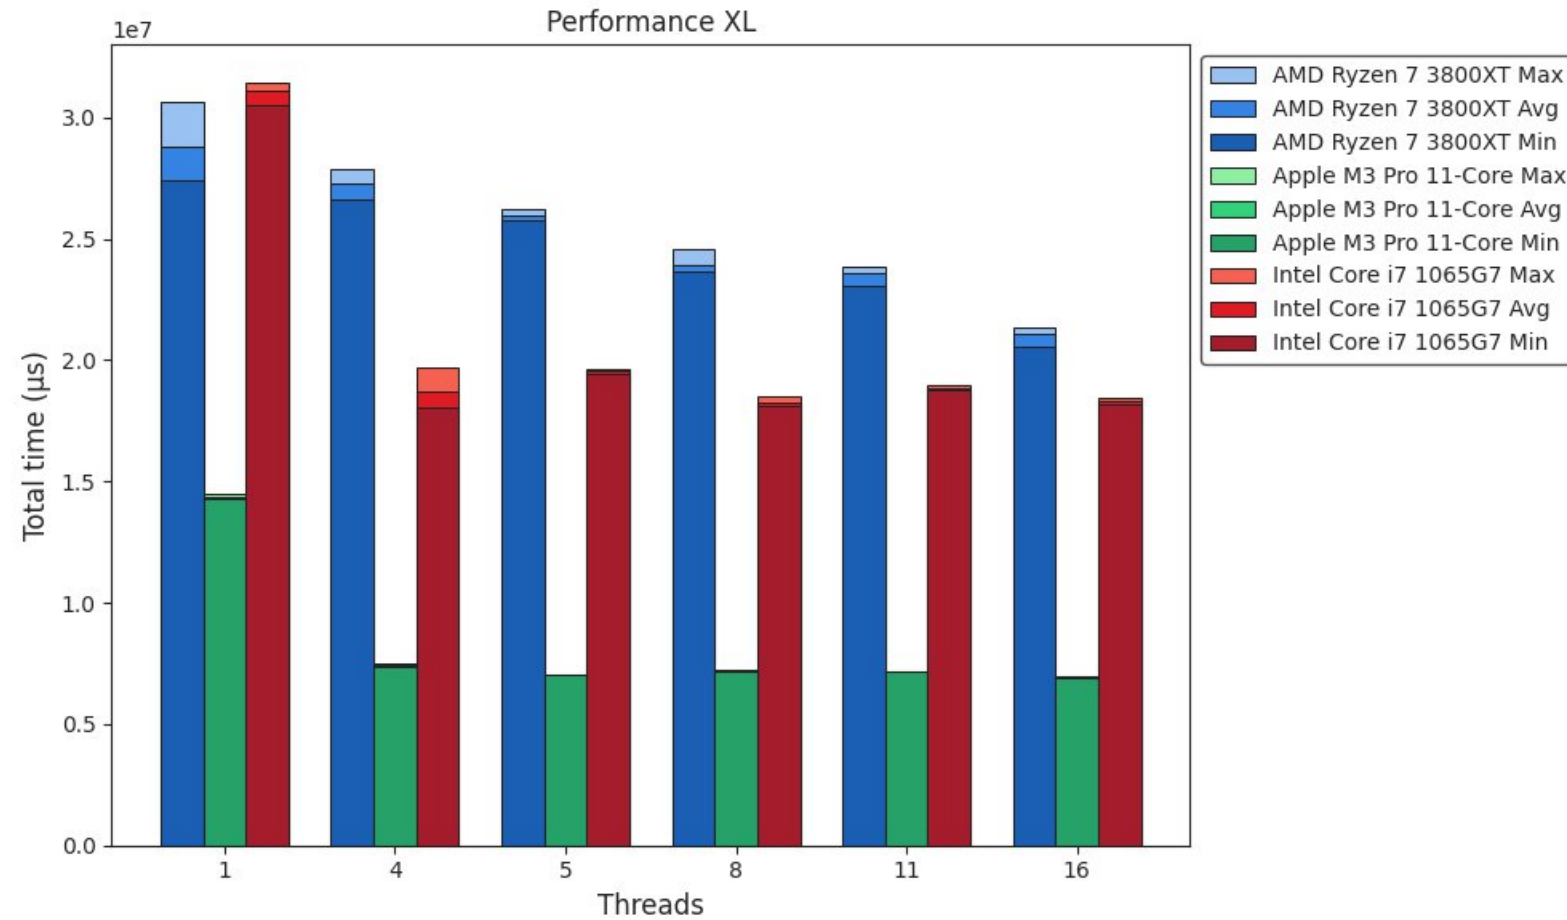
## Performance





# Benchmark

## Performance XL



### ICX:

- Intel oneAPI DPC++/C++ Compiler
- Part of Intel oneAPI
- Formerly Intel C Compiler (ICC)
- Available for Windows and Linux
  - macOS, thus Apple, is not supported
  - → could not be benchmarked
- Supports threading via Intel oneAPI Threading Building Blocks, OpenMP, and native threads

### GCC:


- GNU Compiler Collection
- Available for Windows, Linux, macOS, and a lot more
- Supports threading via OpenMP, native threads, and OpenACC

**Both:** -O3 optimization flag → turns on all optimizations

(We tested the flag for accuracy and it passed)

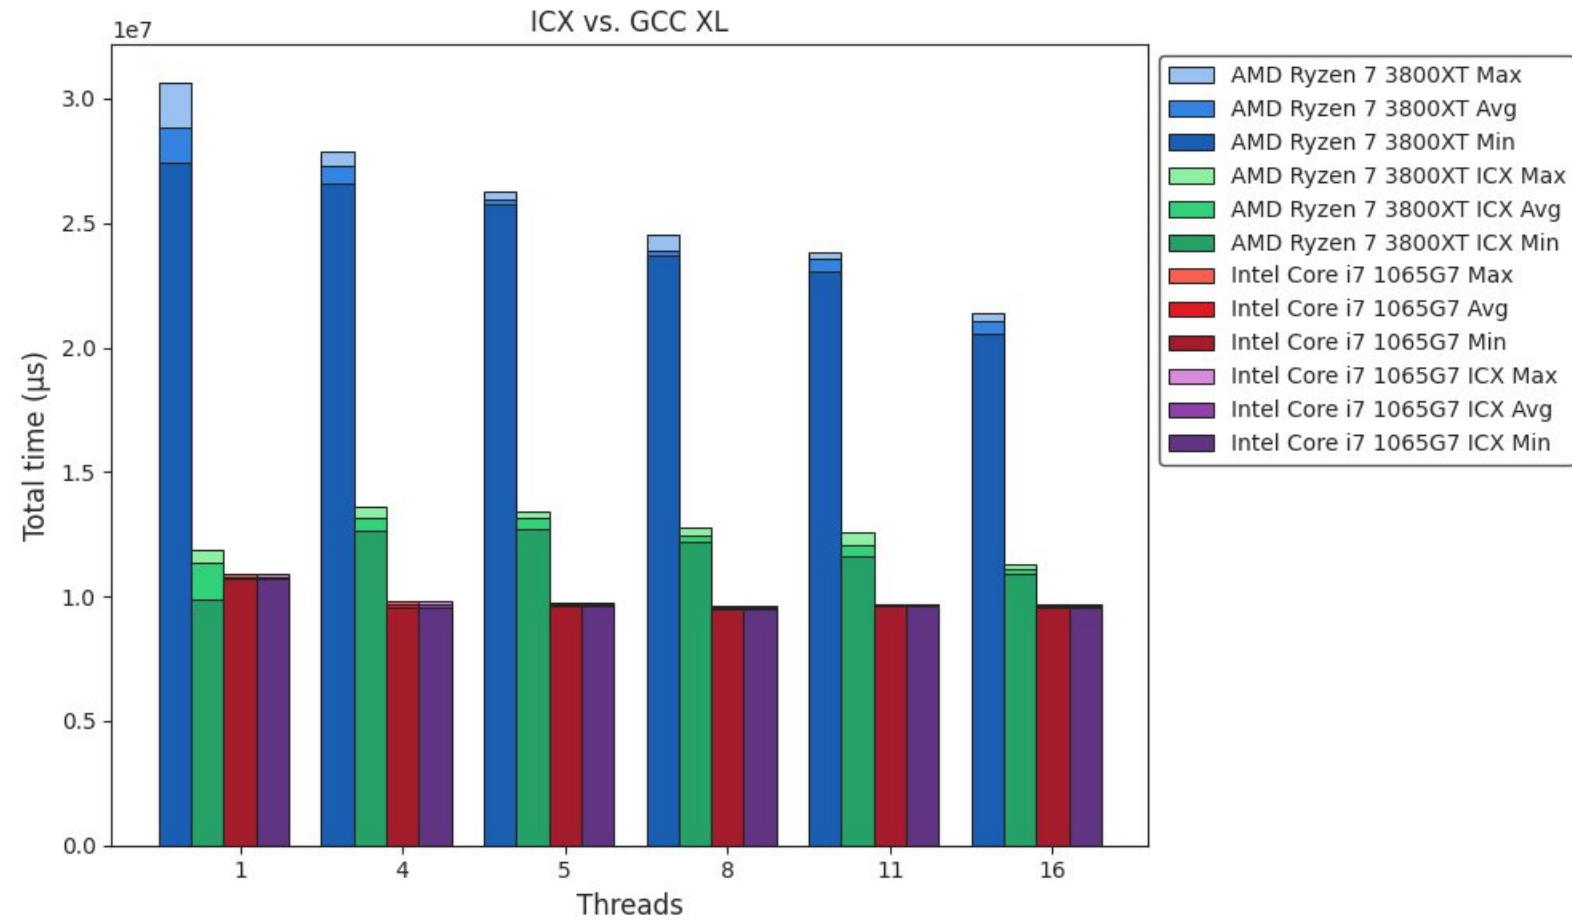
# Benchmark

ICX vs. GCC



# Benchmark

ICX vs. GCC XL



- Implemented matrix multiplication with CUDA in a sandboxed environment
- **First insights:**
  - For our problem size (MNIST), GPU computation takes significantly longer than CPU
    - 206.417 $\mu$ s vs. 2.282 $\mu$ s
    - Reason: Data transfer overhead between CPU and GPU memory dominates computation time
- **Difficulties encountered:**
  - Compiling CUDA C++ correctly within our framework, which is written in C
    - The extern "C" declaration in the .c files is not functioning as expected
    - Using gcc to compile the .c files
    - Using nvcc to compile the .cu files
    - Linking all components together with nvcc



- 
- Framework ✓
  - Multithreading ✓
    - (Implementing OpenMP and comparing its performance to the current solution)
  - ICX vs. GCC ✓
  - SIMD
    - Arm Neon
    - SSE vs. AVX2 vs. AVX-512
  - Quantization
  - CUDA tuning
    - Matrix multiplication ✓
  - (Apple M3 Pro NPU)