

1. Introduction to Deep Learning

Holzinger Philipp



1. Introduction to Deep Learning

1. What is TensorFlow?
2. Building a Graph
3. Basic Training and Inference
 1. Neural Network Fundamentals
 2. The MNIST Dataset
 3. Using Fully Connected Layers
4. Convolutional Neural Networks (CNN)
 1. CNN Motivation
 2. Convolution Basics
 3. Using Convolutional Layers
 4. LeNet

1.1 What is TensorFlow?

1.1 What is TensorFlow?

“TensorFlow™ is an open source software library for high performance numerical computation. Its flexible architecture allows easy deployment of computation across a variety of platforms (CPUs, GPUs, TPUs), and from desktops to clusters of servers to mobile and edge devices. Originally developed by researchers and engineers from the Google Brain team within Google’s AI organization, it comes with strong support for machine learning and deep learning and the flexible numerical computation core is used across many other scientific domains.”

- <https://www.tensorflow.org/>

1.1 What is TensorFlow?

What is a tensor?

A **tensor** is an arbitrarily complex geometric object that maps in a (multi-)linear manner geometric [...] tensors to a resulting tensor. Thereby, vectors and scalars themselves [...] are considered as the simplest tensors.

- <https://en.wikipedia.org/wiki/Tensor>

$$a$$

Scalar
Order 0 tensor

$$\begin{pmatrix} a \\ b \\ c \end{pmatrix}$$

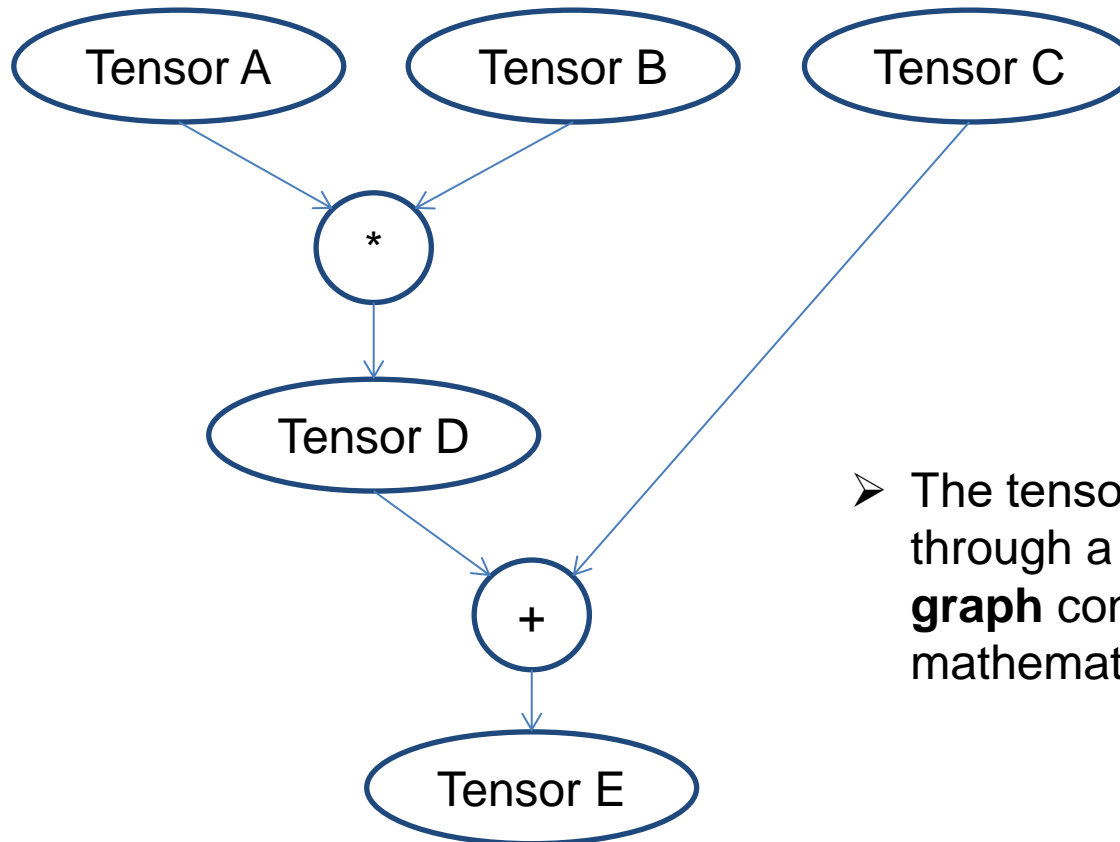
Vector
Order 1 tensor

$$\begin{pmatrix} a & b & c \\ d & e & f \\ g & h & i \end{pmatrix}$$

Matrix
Order 2 tensor

1.1 What is TensorFlow?

What is flowing?



- The tensors are “flowing” through a **computational graph** containing a series of mathematical operations

1.2 Building a Graph

Note:

- Python 3 is used as programming language
- To better understand the fundamental principals of deep neural networks we use the internal low-level TensorFlow 1.x API in this presentation which is closer to the actual C++ code
 - In later versions these functions have been moved to the *tf.compat.v1* namespace (we omit this to shorten the code)
- In practice **always use the high-level APIs like “Keras” to build and train deep learning models**, which also became the default method since TensorFlow 2.x
 - An introduction to Keras can be found here:
https://keras.io/guides/functional_api/

1.2 Building a Graph

Differences between imperative languages and the TensorFlow concept:

- Execution in TensorFlow is graph based
- The TensorFlow graph represents a series of operations to perform
- Typical TensorFlow operations like add, multiply or matmul modify the current graph and do NOT directly perform the operation

Therefore the TensorFlow workflow looks like this:

1. Build a TensorFlow graph with all needed operations
2. Feed the desired input values into the graph
3. Run the graph to get the result

1.2 Building a Graph

```
import tensorflow as tf


# these are tensors of a graph
a = tf.placeholder(tf.float32, shape=(None), name='A')
b = tf.placeholder(tf.float32, shape=(None), name='B')
c = tf.placeholder(tf.float32, shape=(None), name='C')
d = tf.multiply(a,b, name='mul')
e = tf.add(c,d, name='add')

# open execution session
sess = tf.Session()

# this feeds tensors to the placeholders, runs the graph
# and stores the result (= last tensor)
result = sess.run(e, feed_dict={a: [1.2, 2.4], b: [1.0, 0.5],
c: [0.1, 0.2]})

print("the type of result is {}".format(type(result)))
print("the result value of the graph is {}".format(result))
```


TensorBoard graph:



1.3 Basic Training and Inference

1.3.1 Neural Network Fundamentals

When does a single artificial neuron y fire?



x = input
 w = weights
 y = output


$$(w_0 \quad w_1 \quad w_2 \quad w_3) \cdot \begin{pmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{pmatrix} = y_0$$

$$\triangleright \sum_{i=0}^{N-1} w_i \cdot x_i = y$$

$$\triangleright w \cdot x = y \quad (\text{scalar product})$$

1.3.1 Neural Network Fundamentals

When does a single artificial neuron y fire?




The value of y can be additionally modified by an “activation function” f

x = input
 w = weights
 y = output
 $f(\Sigma)$ = activation function

$$\triangleright f\left(\sum_{i=0}^{N-1} w_i \cdot x_i\right) = y$$

$$\triangleright f(w \cdot x) = y$$

What is the activation of multiple artificial neurons?




x = input
 w = weights
 y = output

$$\begin{pmatrix} w_{0,0} & w_{0,1} & w_{0,2} & w_{0,3} \\ w_{1,0} & w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,0} & w_{2,1} & w_{2,2} & w_{2,3} \end{pmatrix} \begin{pmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} y_0 \\ y_1 \\ y_2 \end{pmatrix}$$

- $\sum_{i=0}^{N-1} w_{m,i} \cdot x_i = y_m$
- $Wx = y$ (matrix-vector product)

1.3.1 Neural Network Fundamentals

What is the activation of successive impulses?



$$\begin{pmatrix} w_{0,0} & \cdots & w_{0,3} \\ \vdots & \ddots & \vdots \\ w_{2,0} & \cdots & w_{2,3} \end{pmatrix} \begin{pmatrix} x_{0,0} & x_{0,1} \\ \vdots & \vdots \\ x_{3,0} & x_{3,1} \end{pmatrix} = \begin{pmatrix} y_{0,0} & y_{0,1} \\ y_{1,0} & y_{1,1} \\ y_{2,0} & y_{2,1} \end{pmatrix}$$

x = input
 w = weights
 y = output

$$\triangleright \sum_{i=0}^N w_{m,i} \cdot x_{i,b} = y_{m,b}$$

$$\triangleright WX = Y \quad (\text{matrix-matrix product})$$

Final matrices which describe a “fully connected layer”:

$$\begin{pmatrix} w_{0,0} & \cdots & w_{0,N-1} \\ \vdots & \ddots & \vdots \\ w_{M-1,0} & \cdots & w_{M-1,N-1} \end{pmatrix} \begin{pmatrix} x_{0,0} & \cdots & x_{0,B-1} \\ \vdots & \ddots & \vdots \\ x_{N-1,0} & \cdots & x_{N-1,B-1} \end{pmatrix} = \begin{pmatrix} y_{0,0} & \cdots & y_{0,B-1} \\ \vdots & \ddots & \vdots \\ y_{M-1,0} & \cdots & y_{M-1,B-1} \end{pmatrix}$$

x = input

w = weight

y = output

N = number of inputs

M = number of outputs

B = batch size


1.3.2 The MNIST Dataset

MNIST (Modified National Institute of Standards and Technology) **dataset** is a collection of handwritten digits, commonly used for training neural networks. It contains 60,000 training images and 10,000 testing images



1.3.2 The MNIST Dataset

MNIST (Modified National Institute of Standards and Technology) **dataset** is a collection of handwritten digits, commonly used for training neural networks. It contains 60,000 training images and 10,000 testing images



Some numbers are even hard for humans to classify...

1.3.3 Using Fully Connected Layers

Using a fully connected layer to build the first graph for an MNIST image:

- All MNIST Images have a resolution of 28×28 pixels
 - Input layer has $28 \cdot 28 = 784$ nodes
 - $N = 784$
- They represent a number between 0 and 9
 - Output layer has 10 nodes
 - $M = 10$
- First try connect all of them with one layer
 - $784 \cdot 10 = 7840$ float weights
 - $7840 \cdot 4 = 31360$ bytes of storage

1.3.3 Using Fully Connected Layers

Import the MNIST data and iterate through it:

```
import tensorflow as tf
from official.mnist import dataset

# download dataset and store reference in variable
ds_download_dir = "mnist_dataset"
test_ds = dataset.test(ds_download_dir) # different for every data set

# split data in batches
batch_size = 128
test_ds = test_ds.batch(batch_size)

# define iterator over batches of data
data_iterator =
tf.data.Iterator.from_structure(tf.data.get_output_types(test_ds),tf.data.get_output_shapes(test_ds))

# define graph operation which initializes the iterator with the dataset
test_init_op = data_iterator.make_initializer(test_ds)

# define graph operation which gets the next batch of the iterator over the dataset
next_data_batch = data_iterator.get_next()
```

1.3.3 Using Fully Connected Layers

Building the evaluation graph for an MNIST image:

```
# define initialization parameters
```

```
mu = 0
```

```
sigma = 0.1
```

```
# build evaluation graph with one fully connected layer
```

```
images = tf.placeholder(tf.float32, shape=(None,784),name='images')
```

```
weights = tf.Variable(tf.random.truncated_normal(shape=(784,10), mean=mu, stddev=sigma))
```

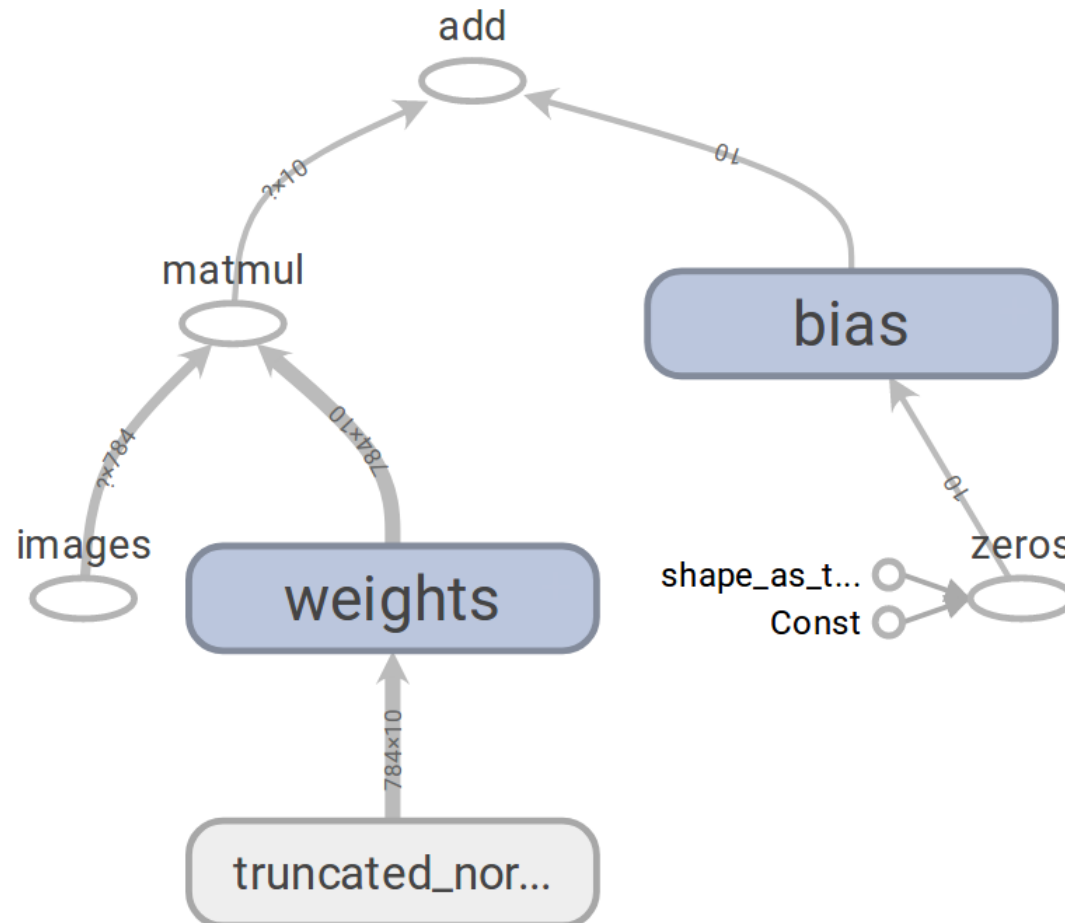
```
bias = tf.Variable(tf.zeros(10), name='bias')
```

```
input_layer = images
```

```
output_layer = tf.add(tf.matmul(input_layer, weights, name='matmul'), bias, name='add')
```

1.3.3 Using Fully Connected Layers

The described graph visualized with TensorBoard:



1.3.3 Using Fully Connected Layers

Building the graph to calculate the accuracy of a batch:

```
# feed the correct labels into the net
labels = tf.placeholder(tf.int32, (None), name='labels')

# highest value is the guess of the network
network_prediction = tf.argmax(output_layer,axis=1,output_type=tf.int32)

# return 0.0 if net prediction is wrong, 1.0 if true
is_correct_prediction = tf.equal(network_prediction, labels)

# percentage of correct predictions is the mean of the batch
accuracy_op = tf.reduce_mean(tf.cast(is_correct_prediction, tf.float32))
```

1.3.3 Using Fully Connected Layers

Running the graphs:

```
# initialize all variables before evaluating the graph
sess = tf.Session()
sess.run(tf.global_variables_initializer())

# initialize the batch iterator
sess.run(test_init_op)

# loop over all batches and calculate predictions and accuracy
while True:
    try:
        data_batch = sess.run(next_data_batch)
        image_batch = data_batch[0]
        label_batch = data_batch[1]
        logits = sess.run(output_layer, feed_dict={images:image_batch})
        accuracy = sess.run(accuracy_op, feed_dict={images:image_batch, labels:label_batch})
    except tf.errors.OutOfRangeError:
        break
```


1.3.3 Using Fully Connected Layers

Accuracy of the graph in this state:

- accuracy: 0.104
- That means 10.4% of all guesses are correct
- 10 numbers must be distinguished (chance of $\frac{1}{10} = 10\%$ to randomly guess correctly)
- Network not much better than random values!
- The weights need to be trained first!

1.3.3 Using Fully Connected Layers

Defining the loss:

```
# define the loss graph
onehot_labels = tf.one_hot(labels, 10)
loss = tf.losses.mean_squared_error(labels=onehot_labels, predictions=output_layer)
```

Loss is the measuring unit of how “wrong” the results really are, e.g.:

Mean Squared Error (MSE):

$$MSE = \frac{1}{M} \sum_{i=0}^{M-1} (c_i - p_i)^2$$

➤ Useful for regression problems

Cross entropy (CE):

$$CE = - \sum_{i=0}^{M-1} c_i \cdot \log(p_i)$$

➤ Useful for classification problems


c_i = correct value (label) p_i = prediction of the network

1.3.3 Using Fully Connected Layers

Minimize the loss:

```
# define training parameters
learning_rate = 0.01

# define the training graph
optimizer = tf.train.GradientDescentOptimizer(learning_rate)
train_op = optimizer.minimize(loss)
```



Update of the weights:

$$W_{t+1} = W_t - \delta E_n X^T$$

δ = learning rate
 E_n = error tensor

1.3.3 Using Fully Connected Layers

Running the graphs (again):

```
# download and prepare training data as before
train_ds = dataset.train(ds_download_dir) # different for every data set
train_ds = train_ds.shuffle(60000).batch(batch_size)
train_init_op = data_iterator.make_initializer(train_ds)

# train the weights by looping repeatedly over all the data (and shuffling in between)
for i in range(epochs):
    sess.run(train_init_op)
    while True:
        try:
            data_batch = sess.run(next_data_batch)
            image_batch = data_batch[0]
            label_batch = data_batch[1]
            sess.run(train_op, feed_dict={images:image_batch, labels:label_batch})
        except tf.errors.OutOfRangeError:
            break

# every epoch calculate predictions, accuracy and loss as before to check progress
<...>
```

1.3.3 Using Fully Connected Layers

Progress during training (MSE loss):

Epoch 0 done: accuracy 0.369, loss 0.173
 Epoch 1 done: accuracy 0.509, loss 0.126
 Epoch 2 done: accuracy 0.581, loss 0.106
 Epoch 3 done: accuracy 0.627, loss 0.095
 Epoch 4 done: accuracy 0.658, loss 0.087
 Epoch 5 done: accuracy 0.680, loss 0.082
 Epoch 6 done: accuracy 0.701, loss 0.078
 Epoch 7 done: accuracy 0.715, loss 0.074
 Epoch 8 done: accuracy 0.726, loss 0.072
 Epoch 9 done: accuracy 0.735, loss 0.069
 Epoch 10 done: accuracy 0.746, loss 0.068
 Epoch 11 done: accuracy 0.752, loss 0.066
 ⋮
 Epoch 96 done: accuracy 0.850, loss 0.042
 Epoch 97 done: accuracy 0.852, loss 0.042
 Epoch 98 done: accuracy 0.851, loss 0.042
 Epoch 99 done: accuracy 0.852, loss 0.042

1.3.3 Using Fully Connected Layers

Accuracy of the graph now:

- Much better, 85.2% of all predictions are correct!

How to further improve the accuracy?


- Use a second fully connected layer
- However: $Y = W_2(W_1X) = (W_2W_1)X = W_{2,1}X$
- Two matrix multiplications in series (two layers) can be expressed as a single one (= one layer)
- Use a suitable activation function $f(\Sigma)$ between the layers!

1.3.3 Using Fully Connected Layers

Examples of activation functions:


Rectified linear unit (ReLU):

$$f(x) = \begin{cases} 0, & \text{when } x < 0 \\ x, & \text{else} \end{cases}$$



Hyperbolic tangent (tanh):

$$f(x) = \tanh(x)$$



1.3.3 Using Fully Connected Layers

Building the evaluation graph for an MNIST image with two layers:

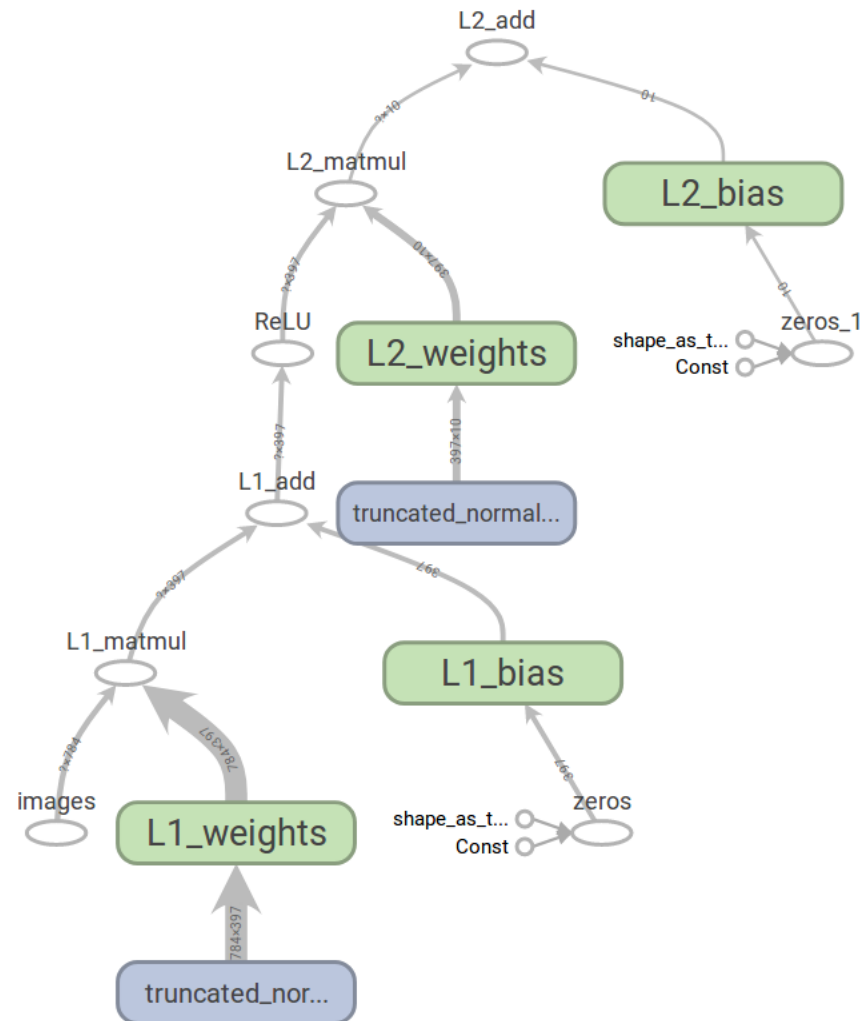
```
# define size of hidden layer
hl_size = 397

# build evaluation graph with two fully connected layers
images = tf.placeholder(tf.float32, shape=(None,784),name='images')
l1_weights = tf.Variable(tf.random.truncated_normal(shape=(784,hl_size), mean=mu, stddev=sigma))
l2_weights = tf.Variable(tf.random.truncated_normal(shape=(hl_size,10), mean=mu, stddev=sigma))
l1_bias = tf.Variable(tf.zeros(hl_size), name='L1_bias')
l2_bias = tf.Variable(tf.zeros(10), name='L2_bias')

input_layer = images
hidden_layer = tf.add(tf.matmul(input_layer, l1_weights, name='L1_matmul'), l1_bias, name='L1_add')
hidden_layer = tf.nn.relu(hidden_layer, name='ReLU')
output_layer = tf.add(tf.matmul(hidden_layer, l2_weights, name='L2_matmul'), l2_bias, name='L2_add')
```


1.3.3 Using Fully Connected Layers

The described graph visualized with TensorBoard:



1.3.3 Using Fully Connected Layers

Progress during training (MSE loss):

Epoch 0 done: **accuracy 0.502**, loss 0.146
 Epoch 1 done: accuracy 0.608, loss 0.110
 Epoch 2 done: accuracy 0.671, loss 0.093
 Epoch 3 done: accuracy 0.713, loss 0.083
 Epoch 4 done: accuracy 0.743, loss 0.075
 Epoch 5 done: accuracy 0.764, loss 0.069
 Epoch 6 done: accuracy 0.781, loss 0.065
 Epoch 7 done: accuracy 0.795, loss 0.061
 Epoch 8 done: accuracy 0.810, loss 0.058
 Epoch 9 done: accuracy 0.820, loss 0.055
 Epoch 10 done: accuracy 0.829, loss 0.053
 Epoch 11 done: accuracy 0.838, loss 0.051
 ⋮
 Epoch 96 done: accuracy 0.930, loss 0.023
 Epoch 97 done: accuracy 0.931, loss 0.022
 Epoch 98 done: accuracy 0.931, loss 0.022
 Epoch 99 done: **accuracy 0.931**, loss 0.022

1.3.3 Using Fully Connected Layers

Progress during training (softmax + CE loss):

Epoch 0 done: accuracy 0.865, loss 0.530
 Epoch 1 done: accuracy 0.893, loss 0.403
 Epoch 2 done: accuracy 0.903, loss 0.355
 Epoch 3 done: accuracy 0.908, loss 0.326
 Epoch 4 done: accuracy 0.915, loss 0.306
 Epoch 5 done: accuracy 0.917, loss 0.289
 Epoch 6 done: accuracy 0.922, loss 0.278
 Epoch 7 done: accuracy 0.924, loss 0.268
 Epoch 8 done: accuracy 0.927, loss 0.258
 Epoch 9 done: accuracy 0.930, loss 0.250
 Epoch 10 done: accuracy 0.932, loss 0.244
 Epoch 11 done: accuracy 0.934, loss 0.236

⋮


Epoch 96 done: accuracy 0.972, loss 0.098
 Epoch 97 done: accuracy 0.972, loss 0.096
 Epoch 98 done: accuracy 0.972, loss 0.097
 Epoch 99 done: accuracy 0.972, loss 0.096

Higher accuracy after
first epoch
➤ Faster training

Slightly better accuracy
after last epoch
➤ Better results


1.4 Convolutional Neural Networks (CNN)

A feature is often only in a subset of the image



How relevant are the pixels in these areas to detect the face?

A feature is often only in a subset of the image



How relevant are the pixels in these areas to detect the face?

How relevant are the pixels in this area to detect the face?

- Pixels which are close together have more common information than pixels far away

1.4.1 CNN Motivation

Using this assumption to improve the weight matrix:

Only weights which correspond to pixels in the proximity are non zero

$$\begin{pmatrix} w_0 \cdots w_i & 0 \cdots 0 & w_j \cdots w_k & \cdots & 0 \cdots 0 & 0 \cdots 0 & 0 \cdots 0 \\ w_l \cdots w_m & 0 \cdots 0 & w_n \cdots w_o & \cdots & 0 \cdots 0 & 0 \cdots 0 & 0 \cdots 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 \cdots 0 & 0 \cdots 0 & 0 \cdots 0 & \cdots & w_p \cdots w_q & 0 \cdots 0 & w_r \cdots w_s \end{pmatrix} \begin{pmatrix} x_0 \\ \vdots \\ x_{N-1} \end{pmatrix} = \begin{pmatrix} y_0 \\ \vdots \\ y_{M-1} \end{pmatrix}$$

Using the many 0 constants to improve memory footprint and runtime:

- Sparse matrix multiplication
- Convolutions

Factors to consider:

- How to choose the size of the neighborhood?
- How many regions are good (= size of M)?

1.4.2 Convolution Basics

Definition Convolution:

Integral over the product of a function with (another) reversed and shifted function:

$$f, g : \mathbb{R} \rightarrow \mathbb{R} \quad (\text{continuous case}):$$


$$(f * g)(n) = \int_{-\infty}^{\infty} f(k) \cdot g(n - k) dk$$

$$f, g : \mathbb{Z} \rightarrow \mathbb{Z} \quad (\text{discrete case}):$$

$$(f * g)[n] = \sum_{k \in \mathbb{Z}} f[k] \cdot g[n - k]$$

1.4.2 Convolution Basics

Image convolved with Gaussian filter to blur edges and reduce noise:



➤ Low pass filter

1.4.2 Convolution Basics

Image convolved with Sobel filter to highlight edges:



$$\begin{pmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{pmatrix}$$

1 X



➤ High pass filter

1.4.2 Convolution Basics

Weight matrix for convolutions (e.g. 3×3):

$$\begin{pmatrix} w_4 & w_5 & 0 & 0 & 0 \cdots 0 & w_7 & w_8 & 0 & 0 & \cdots & 0 & 0 & 0 \\ w_3 & w_4 & w_5 & 0 & 0 \cdots 0 & w_6 & w_7 & w_8 & 0 & \cdots & 0 & 0 & 0 \\ 0 & w_3 & w_4 & w_5 & 0 \cdots 0 & 0 & w_6 & w_7 & w_8 & \cdots & 0 & 0 & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \ddots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & 0 & 0 \cdots 0 & 0 & 0 & 0 & 0 & \cdots & w_3 & w_4 & w_5 \\ 0 & 0 & 0 & 0 & 0 \cdots 0 & 0 & 0 & 0 & 0 & \cdots & 0 & w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_0 \\ \vdots \\ x_{N-1} \end{pmatrix} = \begin{pmatrix} y_0 \\ \vdots \\ y_{M-1} \end{pmatrix}$$

- Use the same filter_height \times filter_width weights as a sliding window for the whole image
- Convolutions are still matrix multiplications
- Convolutions can be more efficiently implemented by stencil operations


1.4.2 Convolution Basics

Convolutions in Tensorflow:

N Input Channels:

N×M Masks:

M Output Channels:



1.4.3 Using Convolutional Layers

Preparing the MNIST dataset (border handling):

```
# reshape datasets to 28 x 28 x 1 pixels (height x width x color channels)
train_ds = train_ds.map(lambda image, label: (tf.reshape(image,[28,28,1]),label))
test_ds = test_ds.map(lambda image, label: (tf.reshape(image,[28,28,1]),label))

# pad images with 1 row/column of pixels on each side for 3 x 3 filter (border handling)
train_ds = train_ds.map(lambda image, label: (tf.pad(image,[[1,1],[1,1],[0,0]],'CONSTANT'),label))
test_ds = test_ds.map(lambda image, label: (tf.pad(image,[[1,1],[1,1],[0,0]],'CONSTANT'),label))

# cache the modified data in memory
train_ds = train_ds.cache()
test_ds = test_ds.cache()

# shuffling and dividing in batches as before
train_ds = train_ds.shuffle(60000).batch(batch_size)
test_ds = test_ds.batch(batch_size)
```

1.4.3 Using Convolutional Layers

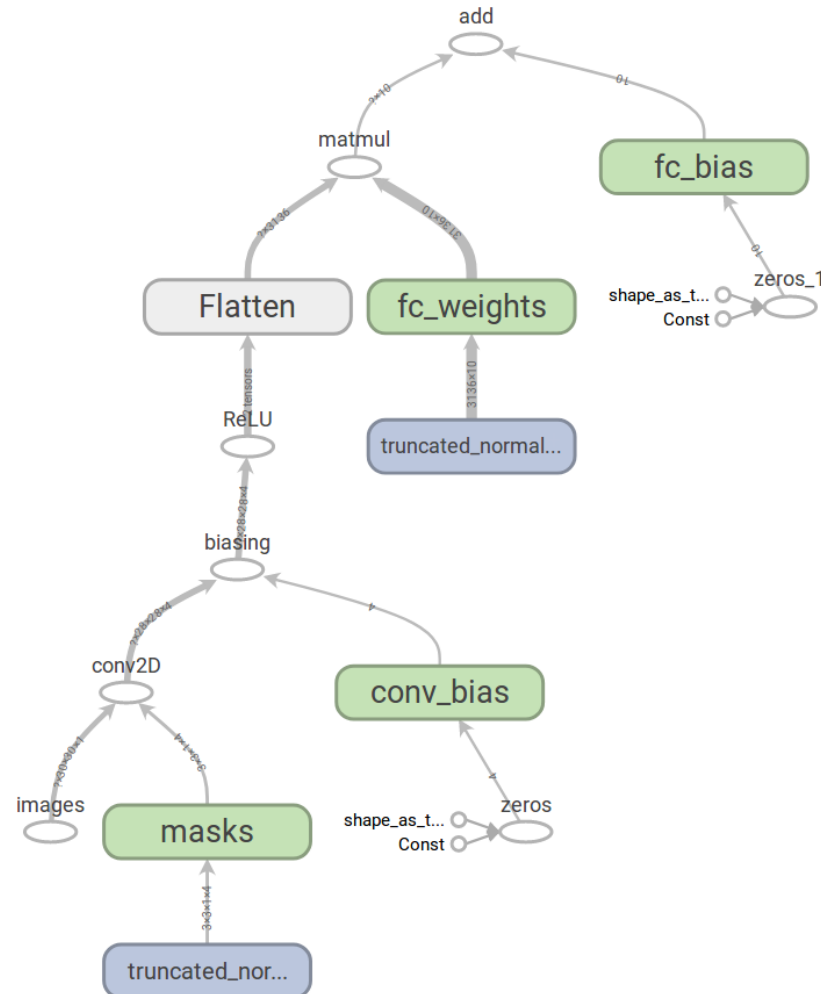
Building the graph with convolution and fully connected layer:

```
# build graph with one convolutional layer (with 4 masks) and one fully connected layer
images = tf.placeholder(tf.float32, shape=(None,30,30,1),name='images')
masks = tf.Variable(tf.random.truncated_normal(shape=(3,3,1,4), mean=mu, stddev=sigma))
fc_weights = tf.Variable(tf.random.truncated_normal(shape=(28*28*4,10), mean=mu, stddev=sigma))
conv_bias = tf.Variable(tf.zeros(4), name="conv_bias")
fc_bias = tf.Variable(tf.zeros(10), name='fc_bias')

input_layer = images
convolution = tf.nn.conv2d(input_layer, masks, strides=[1,1,1,1], padding='VALID', name='conv2D')
convolution = tf.add(convolution, conv_bias, name='biasing')
convolution = tf.nn.relu(convolution, name='ReLU')
hidden_layer = tf.contrib.layers.flatten(convolution)
output_layer = tf.add(tf.matmul(hidden_layer, fc_weights, name='matmul'), fc_bias, name='add')
```

1.4.3 Using Convolutional Layers

The described graph visualized with TensorBoard:



1.4.3 Using Convolutional Layers

Progress during training (softmax + CE loss):

Epoch 0 done: accuracy 0.835, loss 0.552
 Epoch 1 done: accuracy 0.890, loss 0.371
 Epoch 2 done: accuracy 0.903, loss 0.337
 Epoch 3 done: accuracy 0.906, loss 0.317
 Epoch 4 done: accuracy 0.908, loss 0.308
 Epoch 5 done: accuracy 0.913, loss 0.302
 Epoch 6 done: accuracy 0.915, loss 0.296
 Epoch 7 done: accuracy 0.916, loss 0.290
 Epoch 8 done: accuracy 0.918, loss 0.288
 Epoch 9 done: accuracy 0.919, loss 0.281
 Epoch 10 done: accuracy 0.921, loss 0.281
 Epoch 11 done: accuracy 0.922, loss 0.274
 ⋮
 Epoch 96 done: accuracy 0.971, loss 0.095
 Epoch 97 done: accuracy 0.971, loss 0.092
 Epoch 98 done: accuracy 0.971, loss 0.092
 Epoch 99 done: accuracy 0.972, loss 0.090

1.4.3 Using Convolutional Layers

Accuracy of the graph:

- 97.2% of all predictions are correct
- Same accuracy as the example with two fully connected layers


Why convolutions?

- Computations and memory for two fully connected layers:
 - Multiplications: $28 \cdot 28 \cdot 397 + 397 \cdot 10 = 315,218$
 - Additions: $28 \cdot 28 \cdot 397 + 397 \cdot 10 = 315,218$
 - Memory: $(28 \cdot 28 \cdot 397 + 397 \cdot 10) \cdot 4 = 1.26 \text{ MB}$
- Computations and memory for convolution and fully connected layer:
 - Multiplications: $28 \cdot 28 \cdot 3 \cdot 3 \cdot 4 + 28 \cdot 28 \cdot 4 \cdot 10 = 59,584$
 - Additions: $28 \cdot 28 \cdot 3 \cdot 3 \cdot 4 + 28 \cdot 28 \cdot 4 \cdot 10 = 59,584$
 - Memory: $(3 \cdot 3 \cdot 4 + 28 \cdot 28 \cdot 4 \cdot 10) \cdot 4 = 125.58 \text{ KB}$
- 81% less multiplications/additions and 90% less memory!

1.4.3 Using Convolutional Layers

Result of a convolutional layer are M images

- Can be quite large and therefore increases computations
- Are all resulting pixels necessary?
- Pooling as tradeoff between accuracy and computational time



1.4.3 Using Convolutional Layers

Adding maximum pooling to the graph:

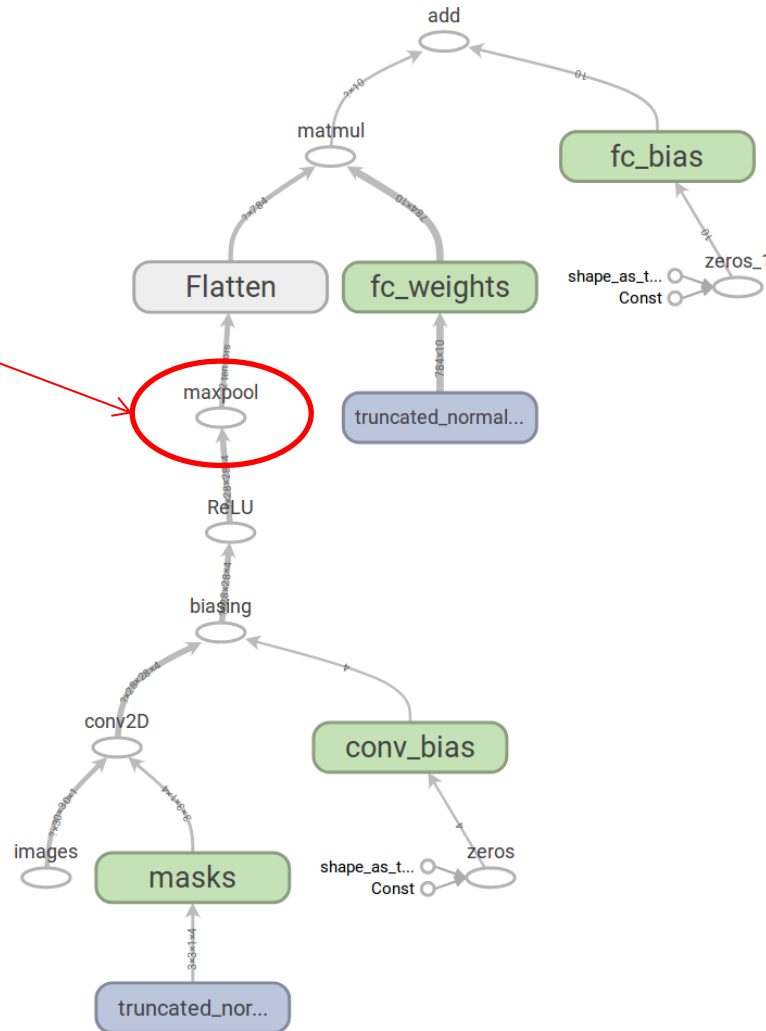
```
# build graph with one convolutional layer (with 4 masks), pooling and one fully connected layer
images = tf.placeholder(tf.float32, shape=(None,30,30,1),name='images')
masks = tf.Variable(tf.random.truncated_normal(shape=(3,3,1,4), mean=mu, stddev=sigma))
fc_weights = tf.Variable(tf.random.truncated_normal(shape=(14*14*4,10), mean=mu, stddev=sigma))
conv_bias = tf.Variable(tf.zeros(4), name="conv_bias")
fc_bias = tf.Variable(tf.zeros(10), name='fc_bias')

input_layer = images
convolution = tf.nn.conv2d(input_layer, masks, strides=[1,1,1,1], padding='VALID', name='conv2D')
convolution = tf.add(convolution, conv_bias, name='biasing')
convolution = tf.nn.relu(convolution, name='ReLU')
pooling = tf.nn.max_pool2d(convolution, ksize=[1,2,2,1], strides=[1,2,2,1], padding='VALID')
hidden_layer = tf.contrib.layers.flatten(pooling)
output_layer = tf.add(tf.matmul(hidden_layer, fc_weights, name='matmul'), fc_bias, name='add')
```

1.4.3 Using Convolutional Layers

The described graph visualized with TensorBoard:

new maximum
pooling layer



1.4.3 Using Convolutional Layers

Progress during training (softmax + CE loss):

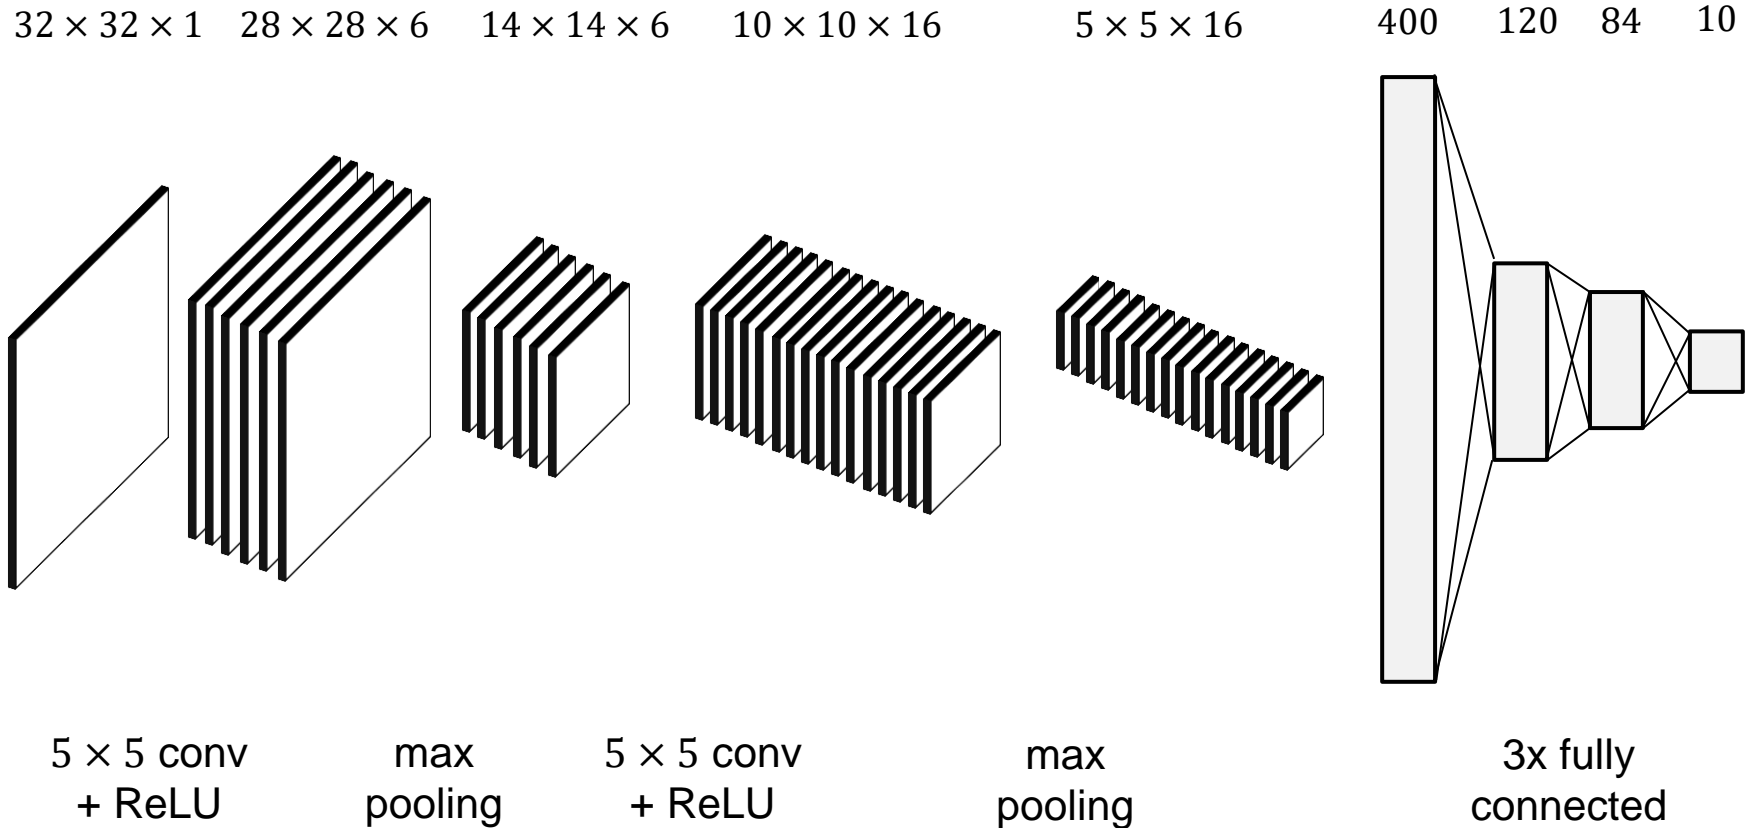
Epoch 0 done: accuracy 0.827, loss 0.579
 Epoch 1 done: accuracy 0.883, loss 0.398
 Epoch 2 done: accuracy 0.901, loss 0.352
 Epoch 3 done: accuracy 0.908, loss 0.329
 Epoch 4 done: accuracy 0.910, loss 0.317
 Epoch 5 done: accuracy 0.912, loss 0.308
 Epoch 6 done: accuracy 0.914, loss 0.301
 Epoch 7 done: accuracy 0.914, loss 0.300
 Epoch 8 done: accuracy 0.915, loss 0.294
 Epoch 9 done: accuracy 0.918, loss 0.290
 Epoch 10 done: accuracy 0.916, loss 0.291
 Epoch 11 done: accuracy 0.918, loss 0.285
 ⋮
 Epoch 96 done: accuracy 0.967, loss 0.106
 Epoch 97 done: accuracy 0.967, loss 0.106
 Epoch 98 done: accuracy 0.967, loss 0.107
 Epoch 99 done: **accuracy 0.968**, loss 0.106

Compared to graph without max pooling:

- **-0.4% points accuracy**
- 39.5% less multiplications/additions
- 75% less memory for weights

1.4.4 LeNet

Yann LeCun's proposed CNN for the MNIST dataset:



- <http://yann.lecun.com/exdb/lenet/>

Progress during training (softmax + CE loss):

Epoch 0 done: accuracy 0.962, loss 0.124
Epoch 1 done: accuracy 0.974, loss 0.080
Epoch 2 done: accuracy 0.984, loss 0.050
Epoch 3 done: accuracy 0.984, loss 0.051
Epoch 4 done: accuracy 0.989, loss 0.033
Epoch 5 done: accuracy 0.987, loss 0.046
Epoch 6 done: accuracy 0.988, loss 0.039
Epoch 7 done: accuracy 0.989, loss 0.035
Epoch 8 done: accuracy 0.988, loss 0.038
Epoch 9 done: accuracy 0.990, loss 0.032
Epoch 10 done: accuracy 0.990, loss 0.036
Epoch 11 done: accuracy 0.989, loss 0.041
⋮
Epoch 96 done: accuracy 0.991, loss 0.073
Epoch 97 done: accuracy 0.991, loss 0.073
Epoch 98 done: accuracy 0.991, loss 0.074
Epoch 99 done: **accuracy 0.991**, loss 0.074